

Lecture 12:

CPU – Superscalar

Hunjun Lee

hunjunlee@hanyang.ac.kr

Our next stuff:

Can we make CPU faster?

◆ The latency of the CPU:

- wall clock time = (time/cyc) x (cyc/inst) x (~~inst/program~~)
 - time/cyc: the critical path delay
 - cyc/inst: the number of cycles to execute a instruction
 - inst/program: ISA and compiler determine this

필요한 수량
→ 무시

How to make a faster CPU?

(1) Reduce time / cyc

~~(2)~~ Increase inst / cyc → IPC ↑

ILP: Instruction-level parallelism

동시에 실행될 수 있는 Instruction의 수

- ◆ ILP is the parallel or simultaneous execution of a sequence of instructions
 - Inter-dependent instructions cannot be executed in parallel
- ◆ Program ILP = Avg. # of instructions / Cycle (step)
 - How many instructions are simultaneously executed in parallel

Data dependency
from 012
실행 순서
→ ILP = 1

code1:

```
addi $r1, $r2, 1
divi $r3, $r1, 17
sub  $r4, $r0, $r3
```

Max ILP = 1
(execute serially)

Data dependency X
→ 실행 순서
실행 가능
→ ILP = 3

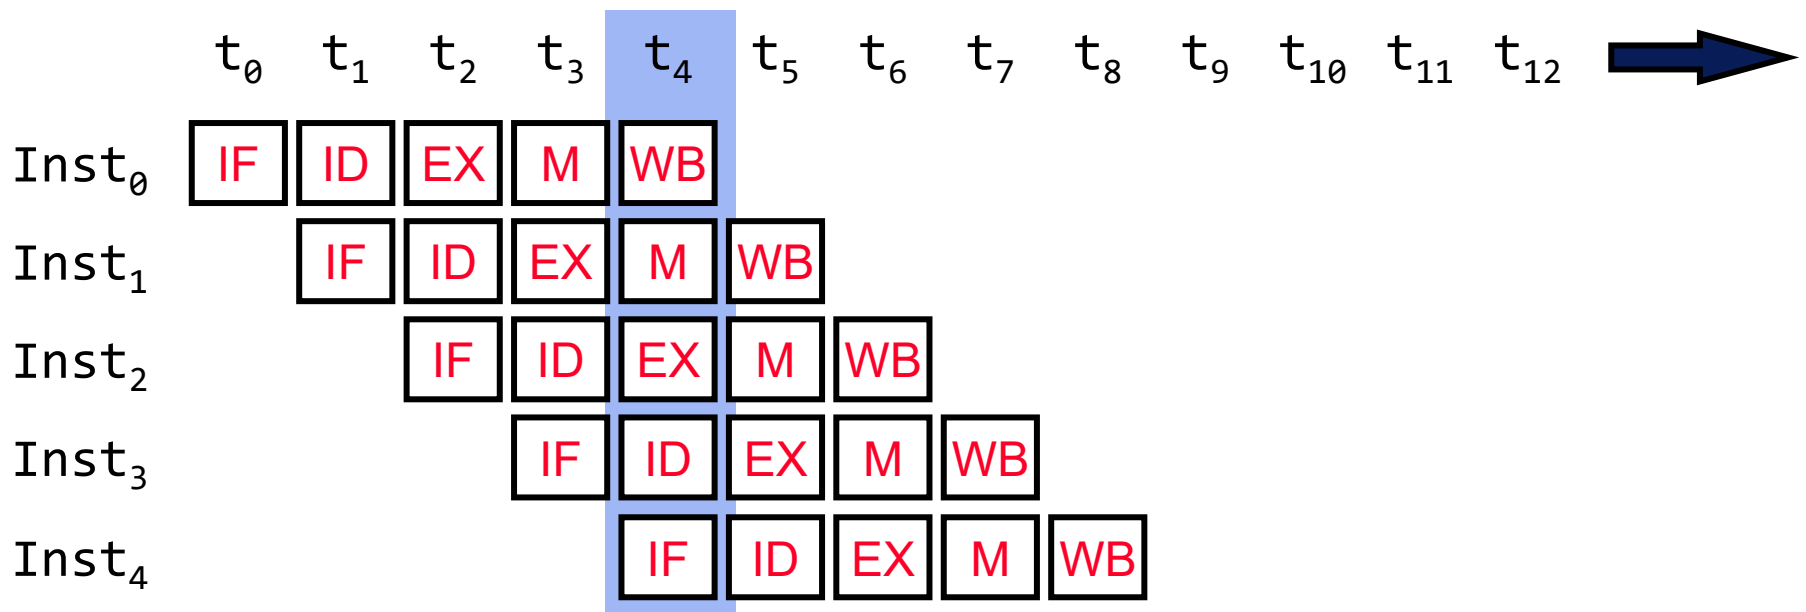
code2:

```
addi $r1, $r2, 1
divi $r3, $r9, 17
sub  $r4, $r0, $r10
```

Max ILP = 3
(execute parallel)

How to exploit ILP: Pipeline


- ◆ Pipelining: executing multiple instructions in parallel
 - **Operation latency = 1**
 - **Peak IPC = 1**
 - **HW ILP** = # of instructions / # of cycles required = **1** $ILP = 1.0 \text{ GHz}$
 - We parallelize instructions at the cost of #cycles



Make the pipeline deeper (superpipeline)

→ 좀더 fetch하는 것 같다.

- ◆ Superpipeline execution (split a cycle into M minor cycles):
 - **Operation latency** = 1 baseline cycle (M minor cycles)
 - **Peak IPC** = M per baseline cycle
 - **HW ILP** = # of instructions / # of (baseline) cycles required = M



This assumes higher
clock frequency...

Recall: Is 5-stage pipeline sufficient?


- ◆ There are still plenty of combinational delay between registers
- ◆ “**Superpipelining**” increases pipelining degree such that even intrinsic operations (e.g., ALU, RF read/write, memory access) require multiple stages
- ◆ Potential problem: “more data hazards”



Pipeline stall even w/ forwarding

Recall: But above all, you cannot pipeline till infinity!

- ◆ There are diminishing returns on clock speed

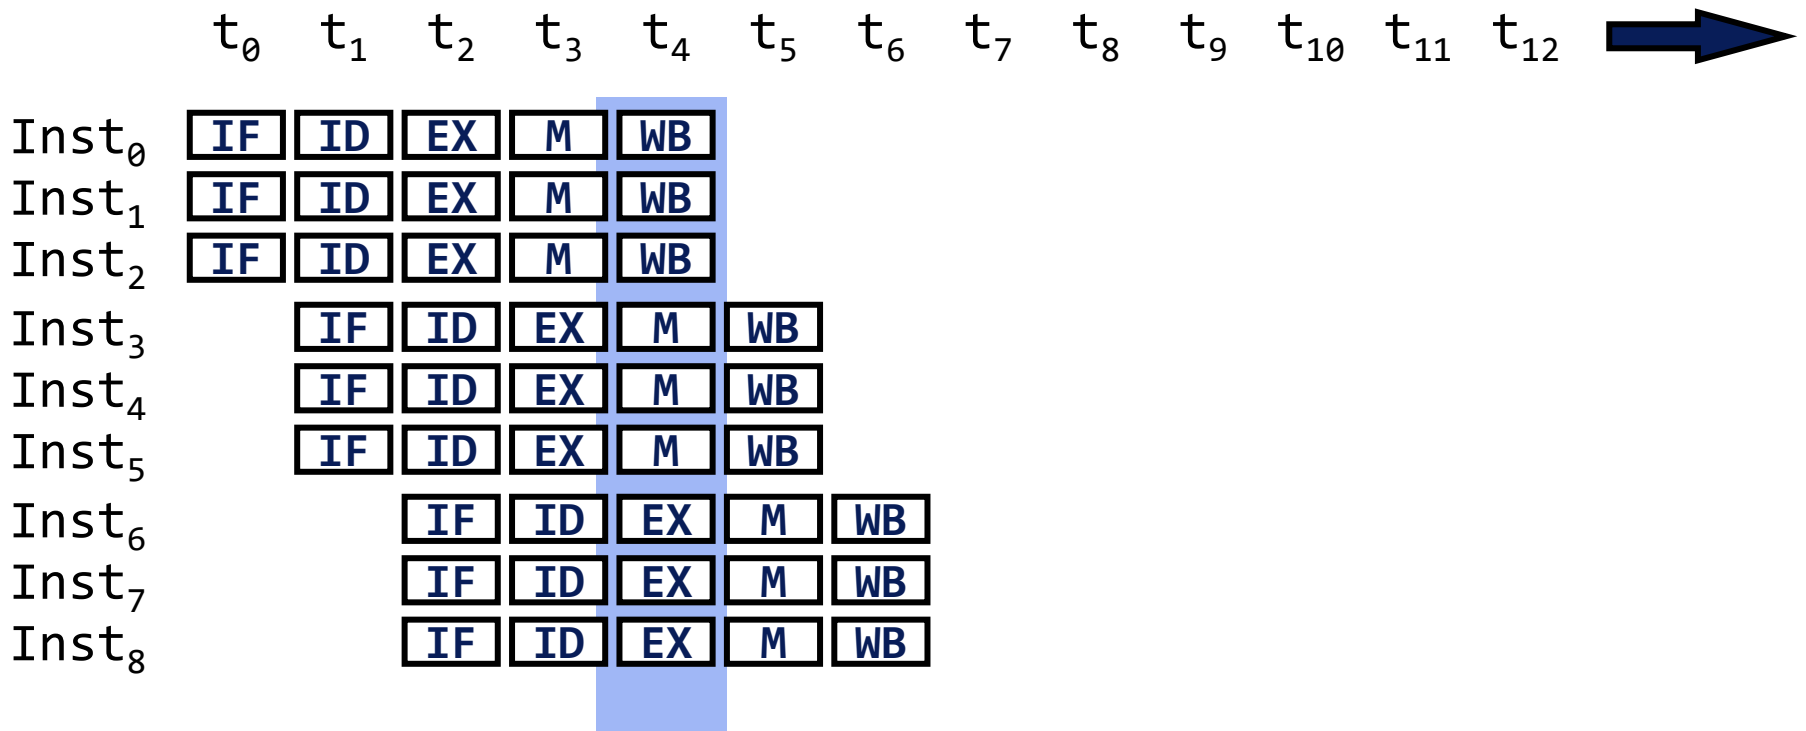


- ◆ Complicates hardware (there should be tons of wires for data forwarding)
- ◆ You cannot always solve the dependency problem w/ forwarding!

We need to use superscalar!!!

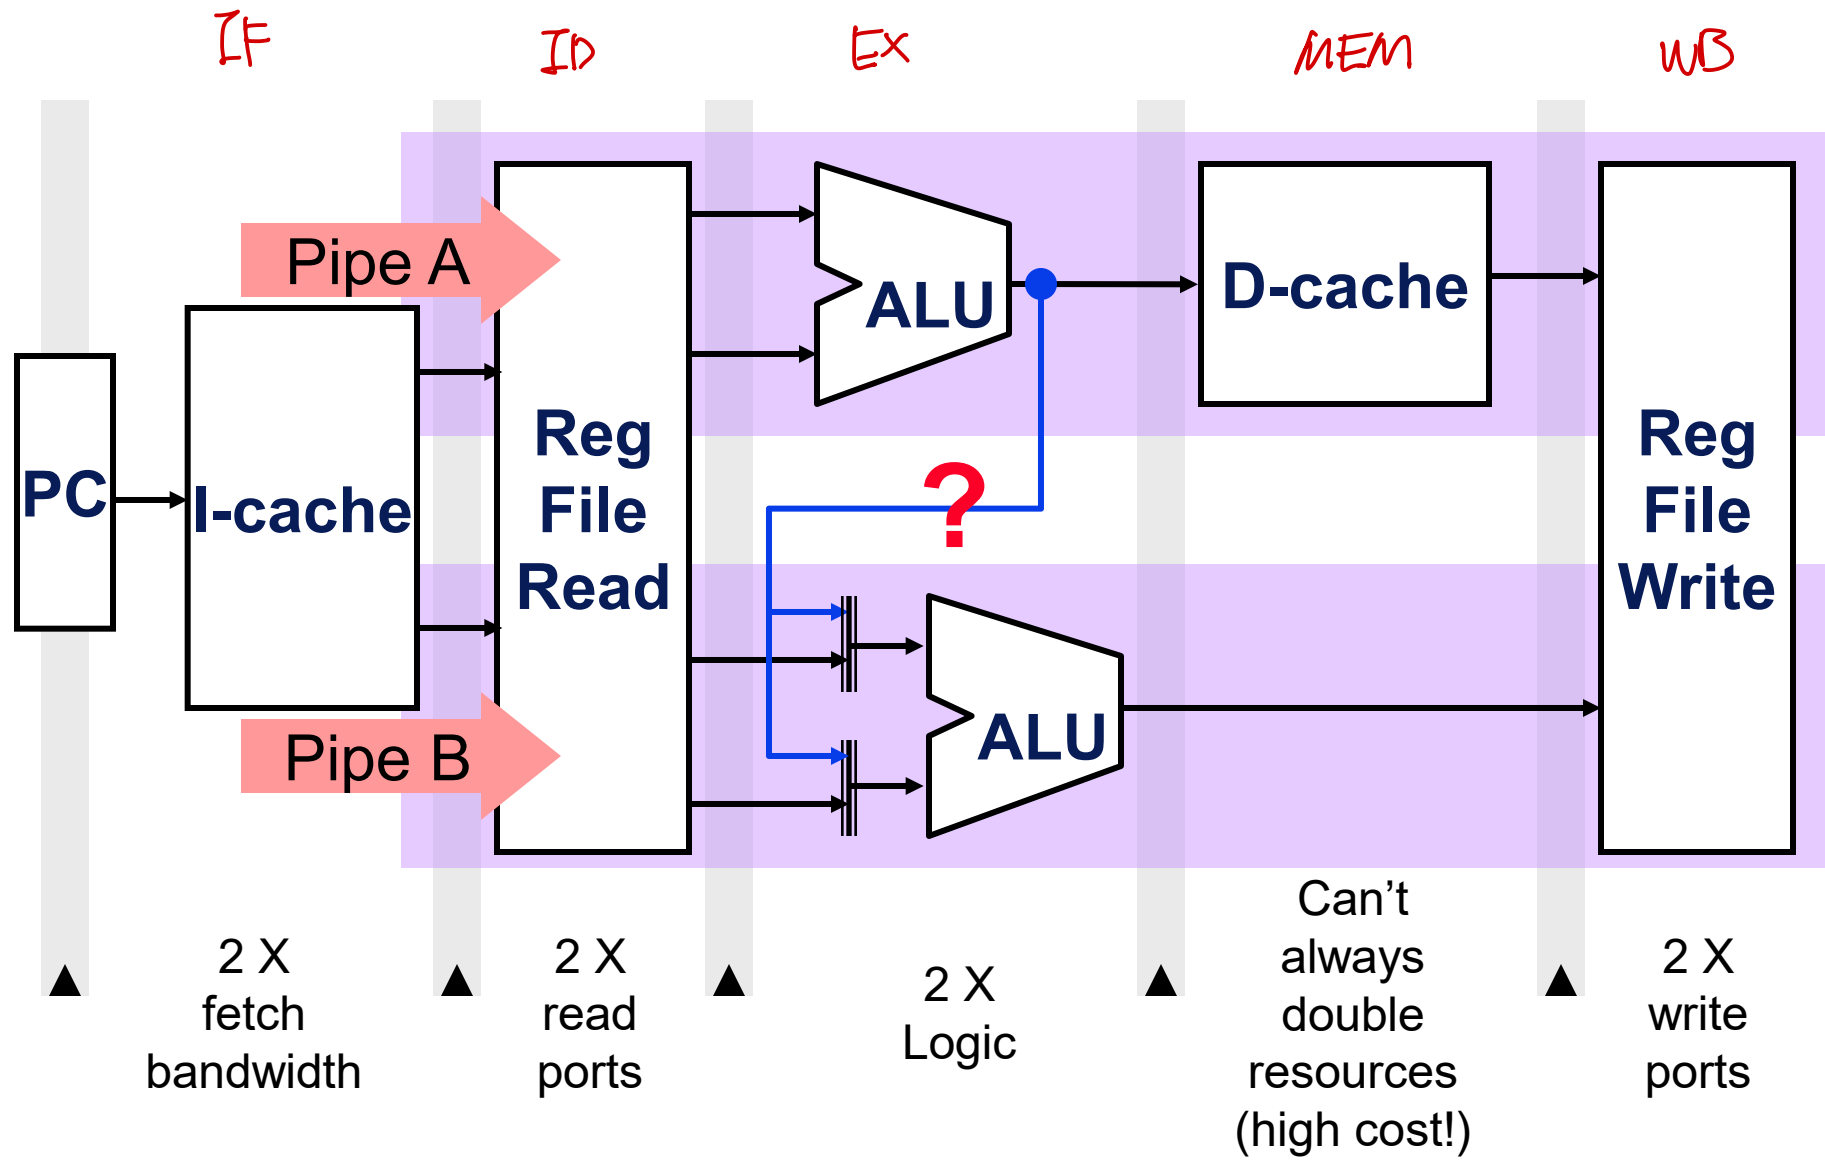
Superscalar Machines

- ◆ Superscalar (+ pipelined) execution
 - Operation latency = 1 baseline cycle Peak throughput
→ N times
 - Peak IPC = N per baseline cycle
 - HW ILP = # of instructions / # of cycles required = N

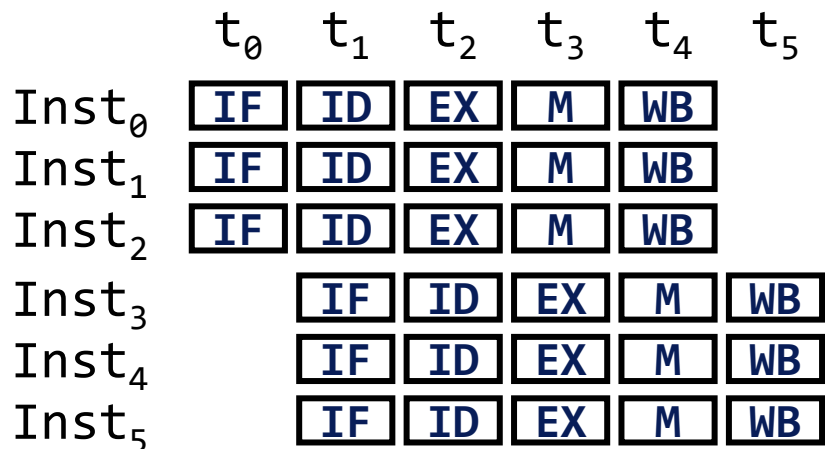


다중 클럭 레지스터 2개
 다중 레지스터 2개 연산 클럭

Example 2-way superscalar datapath



Superscalar vs. Superpipelined




Superscalar Parallelism

Operation Latency: 1

Issuing Rate: **N**

Superscalar Degree: **N**

VS.



Superpipeline Parallelism

Operation Latency: 1

Issuing Rate: **M**

Superpipelined Degree: **M**

We need to exploit both methodologies (N and M) to maximize the performance benefits

How to maximize the ILP?

Instruction scheduling

◆ Static scheduling (or VLIW)

- Compiler groups instructions to be issued together
- Packages them into “issue slots” *동시에 실행 가능 Instruction은*
- Compiler detects and avoids hazards *Instruction을 Grouping 해줌*

◆ Dynamic scheduling

- CPU examines instruction stream and chooses instructions to issue each cycle
- Compiler can help by reordering instructions (out-of-order execution)
- CPU resolves hazards using advanced techniques at runtime

Compiler Cost : $\text{Static} > \text{Dynamic}$
동작

VLIW (Very Long Instruction Word)

◆ Compiler groups instructions into “issue packets”

- It determines the group of instructions that can be issued on a single cycle
- Determined by pipeline resources!

하나의 Cycle에서
dependency, resource 확인

Address가 몇개인지 등

◆ You can make a very long instruction (specifies multiple concurrent operations)

CPU 일종해서는 하나씩 Instruction (packet)을
일어뜨리게, 여러개씩 Instruction을 묶어놓은 것

◆ Compiler must remove some/all hazards

- Reorder instructions into issue packets
- No dependencies within a packet

관찰해 time에 모든 data dependency를
찾아내서 → 성능 ↑

- True or false dependencies (or both?) ← 둘 다 고려!
- Pad with nop if necessary (cannot be parallelized)

- Possibly some dependencies between packets

MIPS with static dual issue

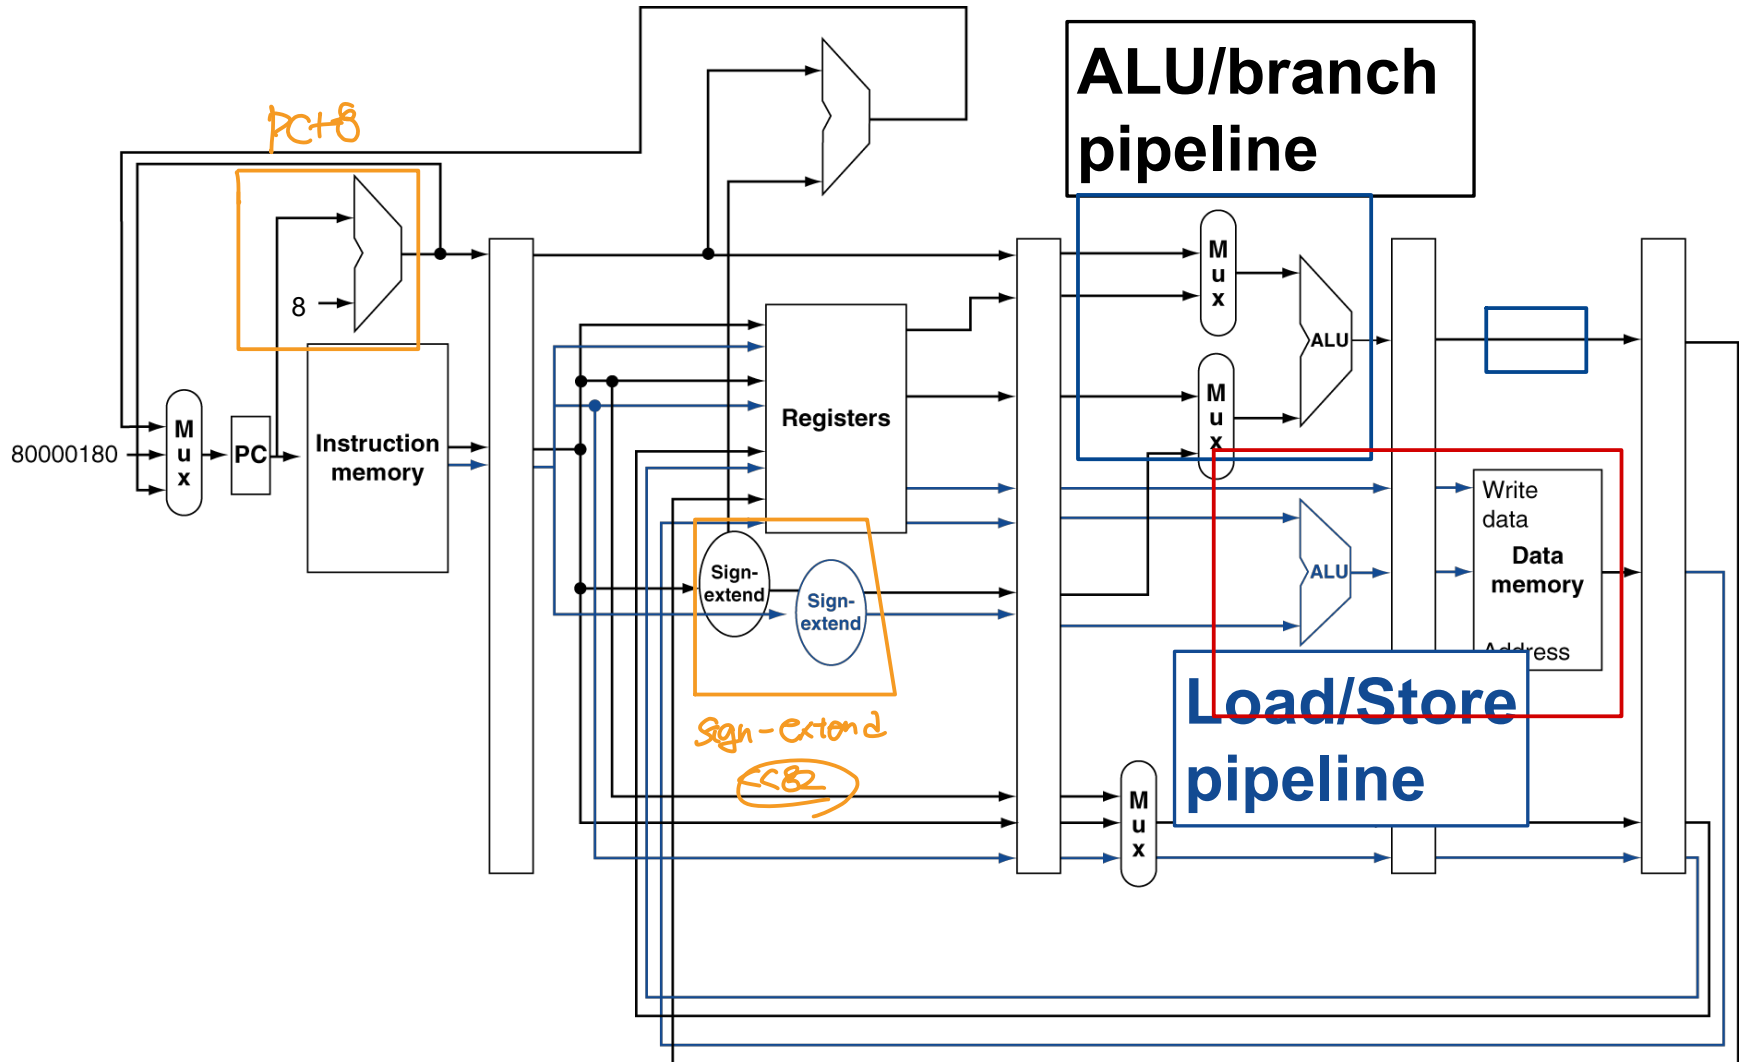
◆ Two-issue packets! : *하루 packet 2개 2개의 Instruction*

- One ALU/branch instruction ← *MEM stage X*
- One load/store instruction
- 64-bit aligned
 - ALU/branch, then load/store
 - Pad an unused instruction with nop ← *노작*

PC 다음 PC+8에 Fetch

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	Bubble MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	Bubble MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	Bubble MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

MIPS with static dual issue



Hazards in the dual-issue MIPS

- ◆ More instructions are executed in parallel
- ◆ **EX data hazard**
 - Can't use ALU result in load/store in same packet


Slot 0 { add **\$t0**, \$s0, \$s1

Slot 1 { load \$s2, 0(**\$t0**)

Split into two packets, effectively a stall

- ◆ **Load-use hazard**

- Still one cycle use latency

Slot 0 { load **\$t0**, 0(\$s0)  *1 cycle stall*

Slot 1 { add \$t2, **\$t0**, \$s1

Split into two packets, effectively a stall

Hazards in the dual-issue MIPS

- ◆ It also suffers from false dependencies

- ◆ **Write after read hazard**

- You cannot place two instructions with WAR hazard

add \$t0, **\$s0**, \$s1

load **\$s0**, 0(\$t1)

Split into two packets, effectively a stall

- ◆ **Write after write hazard**

- The two packed instructions cannot write to the same register

load **\$t0**, 0(\$s0)

add **\$t0**, \$t1, \$s1

Split into two packets, effectively a stall

Scheduling example

- ◆ Schedule instructions in the loop for dual-issue MIPS!

```

addi $s1, $zero, 400           # $s1 = 400
Loop:
    lw    $t0, 0($s1)           # $t0=array element
    addu  $t0, $t0, $s2         # add scalar in $s2
    sw    $t0, 0($s1)           # store result
    addi  $s1, $s1, -4          # decrement pointer
    bne   $s1, $zero, Loop      # branch $s1!=0
  
```

Handwritten note: 100th loop (with arrow pointing to the loop exit branch)

	ALU/branch	Load/store	cycle
Loop:	nop <i>Delay Slot</i>	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1, -4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

Handwritten red arrow: points from the 'Delay Slot' in cycle 1 to the 'addu' instruction in cycle 3.

$$\text{IPC} = 5 / 4 = 1.25 \text{ (2-way superscalar)}$$

Loop unrolling

4번 Unrolling
→ 100번 도는 100P를
25번 돌도록

- ◆ Replicate loop body to expose more parallelism + reduce loop-control overhead

```
# Unroll loop by 4
addi $s1, $zero, 400                                # $s1 = 400
Loop:
    lw    $t0, 0($s1)
    addu  $t0, $t0, $s2
    sw    $t0, 0($s1)
    lw    $t0, -4($s1)
    addu  $t0, $t0, $s2
    sw    $t0, -4($s1)
    lw    $t0, -8($s1)
    addu  $t0, $t0, $s2
    sw    $t0, -8($s1)
    lw    $t0, -12($s1)
    addu  $t0, $t0, $s2
    sw    $t0, -12($s1)
    addi  $s1, $s1, -16
    bne   $s1, $zero, Loop # branch $s1!=0
```

Decrement +
Branch once per
four loops

← 원래 -4, 4 Unrolling이라 -16

Loop unrolling & Scheduling

쓰기만 ↓, Cycle만 ↑
↓ 쓰기 → Low IP Cache

한 Loop에서 명령어 5개, 4개 Cycle 4개 → 4개 Loop에서 명령어 20개, 22개 Cycle 16개

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0 , 0(\$s1)	1
	nop (stall)	nop	2
	addu \$t0 , \$t0 , \$s2	nop	3
	nop	sw \$t0 , 0(\$s1)	4
	nop	lw \$t0 , -4(\$s1)	5
	nop (stall)	nop	6
	addu \$t0 , \$t0 , \$s2	nop	7
	nop	sw \$t0 , -4(\$s1)	8
	nop	lw \$t0 , -8(\$s1)	9
	nop (stall)	nop	10
	addu \$t0 , \$t0 , \$s2	nop	11
	nop	sw \$t0 , -8(\$s1)	12
	nop	lw \$t0 , -12(\$s1)	13
	addi \$s1 , \$s1, -16	nop	14
	addu \$t0 , \$t0 , \$s2	nop	15
	bne \$s1 , \$zero, Loop	sw \$t0 , 4(\$s1)	16

Unrolling

WAW

WAW

WAR

Loop unrolling & Scheduling

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0 , 0(\$s1)	1
	nop (stall)	nop	2
	addu \$t0 , \$t0 , \$s2	nop	3
	nop	sw \$t0 , 0(\$s1)	4
	nop	lw \$t0 , 4(\$s1)	5
	nop (stall)	nop	6
	nop	sw \$t0 , 4(\$s1)	7
	nop	lw \$t0 , 8(\$s1)	8
	nop (stall)	nop	9
	nop	sw \$t0 , 8(\$s1)	10
	nop	lw \$t0 , 12(\$s1)	11
	nop (stall)	nop	12
	addi \$s1 , \$s1, -16	lw \$t0 , -12(\$s1)	13
	addu \$t0 , \$t0 , \$s2	nop	14
	bne \$s1 , \$zero, Loop	nop	15
		sw \$t0 , 4(\$s1)	16

16 cycles per 4 loops (no cycle reduction)
even though the # of instructions are reduced

Because ... consider $IPC = (14 / 16) / < 1.25$
→ We do not benefit from the
improved parallelization opportunities

Why low IPC?

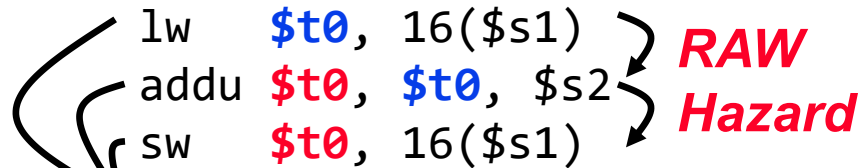
- ◆ There is an **extremely large # of dependencies**

```
# Unroll loop by 4
addi $s1, $zero, 400
```

```
# $s1 = 400
```

```
Loop:
```

```
    lw    $t0, 16($s1)
    addu   $t0, $t0, $s2
    sw     $t0, 16($s1)
```



RAW

Hazard

WAW / WAR
Hazard

```
    lw    $t0, 12($s1)
    addu   $t0, $t0, $s2
    sw     $t0, 12($s1)
    lw    $t0, 16($s1)
    addu   $t0, $t0, $s2
    sw     $t0, 16($s1)
    lw    $t0, 16($s1)
    addu   $t0, $t0, $s2
    sw     $t0, 16($s1)
    addi   $s1, $s1, -16
    bne    $s1, $zero, Loop # branch $s1!=0
```

SW-based Register renaming

- ◆ Compiler renames registers to remove false dependencies

```
# Unroll loop by 4
addi $s1, $zero, 400                # $s1 = 400
Loop:
    lw    $t0, 16($s1)
    addu  $t0, $t0, $s2
    sw    $t0, 16($s1)
    lw    $t1, 12($s1)
    addu  $t1, $t1, $s2
    sw    $t1, 12($s1)
    lw    $t2, 16($s1)
    addu  $t2, $t2, $s2
    sw    $t2, 16($s1)
    lw    $t3, 16($s1)
    addu  $t3, $t3, $s2
    sw    $t3, 16($s1)
    addi  $s1, $s1, -16
    bne   $s1, $zero, Loop # branch $s1!=0
```

Scheduling after renaming

	ALU/branch	Load/store	cycle
Loop:	addi \$s1 , \$s1, -16	nop	1
	nop	lw \$t0 , 16(\$s1)	2
	nop	lw \$t1 , 12(\$s1) <i>stall the lw due to reg</i>	3
	addu \$t0 , \$t0 , \$s2	lw \$t2 , 8(\$s1)	4
	addu \$t1 , \$t1 , \$s2	lw \$t3 , 4(\$s1)	5
	addu \$t2 , \$t2 , \$s2	sw \$t0 , 16(\$s1)	6
	addu \$t3 , \$t3 , \$s2	sw \$t1 , 12(\$s1)	7
	nop	sw \$t2 , 8(\$s1)	8
	bne \$s1 , \$zero, Loop	sw \$t3 , 4(\$s1)	9

Scheduling after renaming

	ALU/branch	Load/store	cycle
Loop:	addi \$s1 , \$s1, -16	nop	1
	nop	lw \$t0 , 16(\$s1)	2
	nop	lw \$t1 , 12(\$s1)	3
	addu \$t0 , \$t0 , \$s2		4
	addi \$s1 , \$s1 , -16		5
			6
			7
			8
			9

9 cycles per 4 loops (7 cycle reduction)

IPC가 증가하는 효과!

Consider $IPC = 1.555 > 1.25 \rightarrow \frac{20}{16} = 1.25$

→ Benefits from the improved parallelization opportunities

Sum up ...

Renaming & scheduling increases ILP!!

- ◆ We can execute previously dependent (false dependency) instructions in parallel



```
div r1 r2 r3
mul r4 r1 r5
add r1 r3 r6
sub r3 r1 r5
```

Original

```
div r1 r2 r3
mul r4 r1 r5
add r8 r3 r6
sub r9 r8 r5
```

Rename

```
div r1 r2 r3
add r8 r3 r6
mul r4 r1 r5
sub r9 r8 r5
```

Rename
+ Reorder

```
{ div r1 r2 r3
  add r8 r3 r6
  { mul r4 r1 r5
    sub r9 r8 r5
```

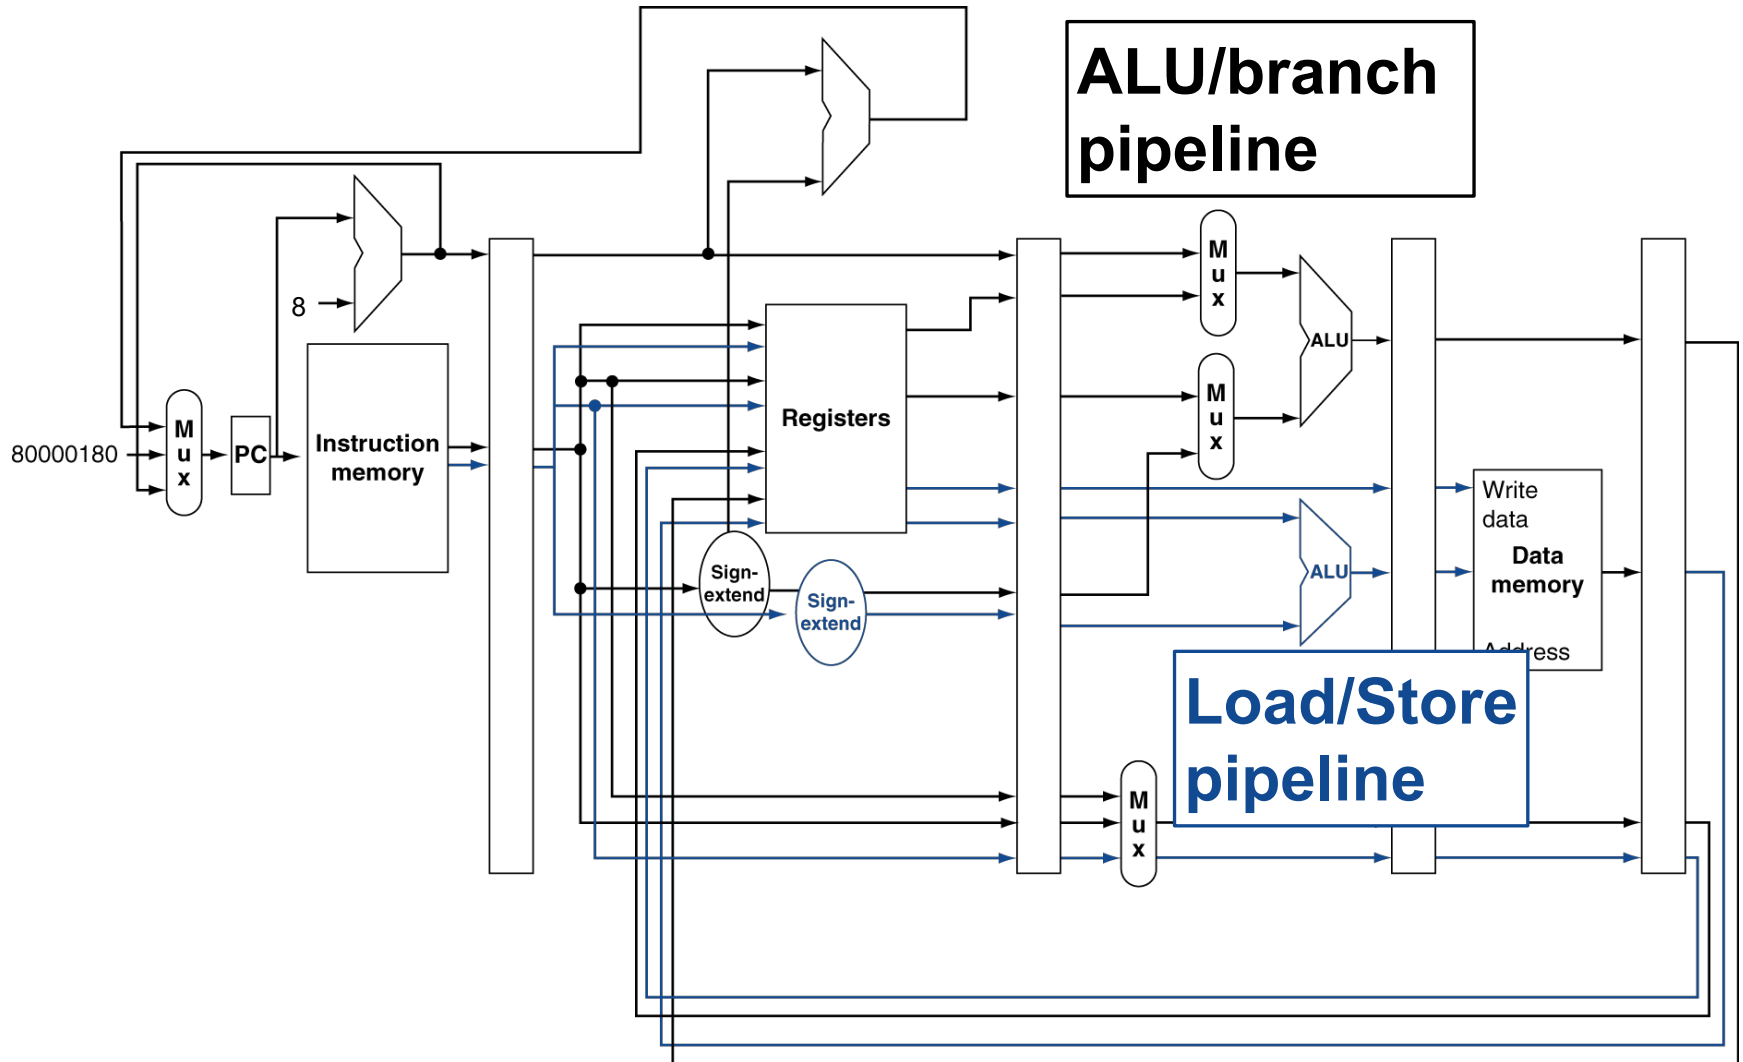
Original
+ Superscalar

MIPS with static dual issue

- ◆ Two-issue packets!
 - One ALU/branch instruction
 - One load/store instruction
 - 64-bit aligned
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

MIPS with static dual issue



How to maximize the ILP?

Instruction scheduling

◆ **Static scheduling (or VLIW)**

- Compiler groups instructions to be issued together
- Packages them into “issue slots”
- Compiler detects and avoids hazards

◆ **Dynamic scheduling**

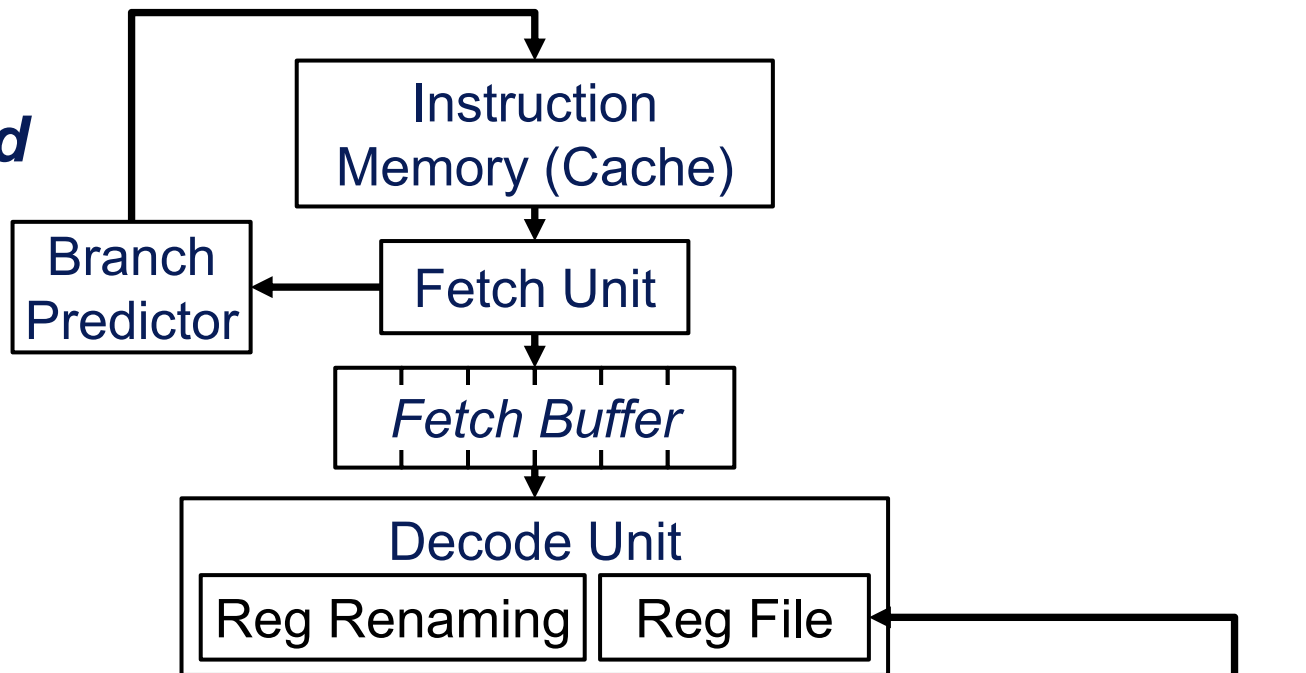
- CPU examines instruction stream and chooses instructions to issue each cycle
- Compiler can help by reordering instructions (out-of-order execution)
- CPU resolves hazards using advanced techniques at runtime

Superscalar Processing

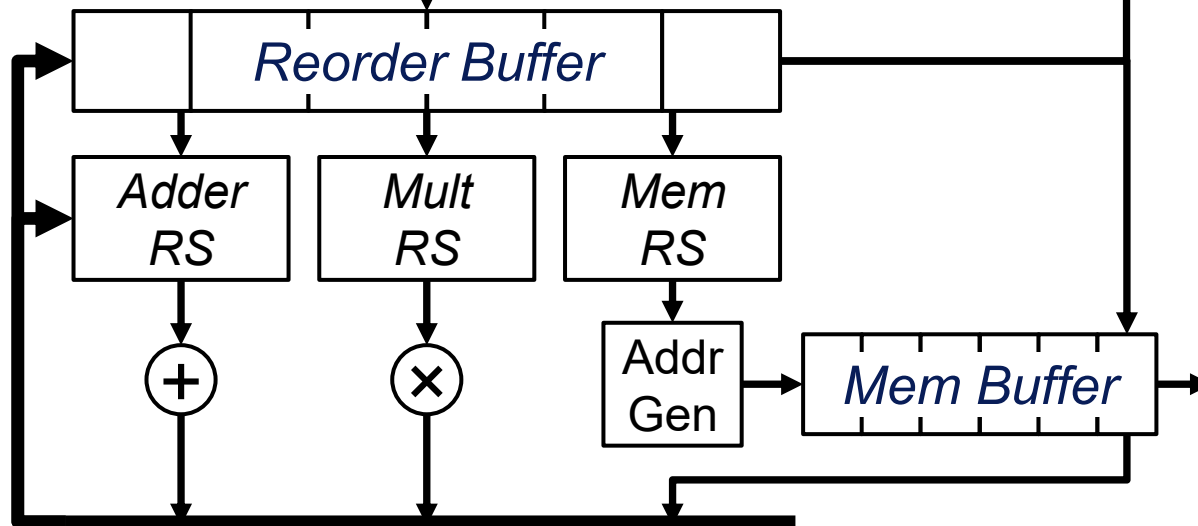
- ◆ Parallel N-instruction fetch
- ◆ Parallel N-instruction decode
- ◆ Parallel N-instruction renaming → (detect dependencies in parallel)
 - Simultaneously allocate ROB, RS ...
- ◆ Dispatch N instructions in parallel (different execution units)
- ◆ Commit N instructions in parallel (N write ports RF)
- ◆ ...

→ 동시에 N개씩 명령어를
인수할 수 있도록 ~

FrontEnd

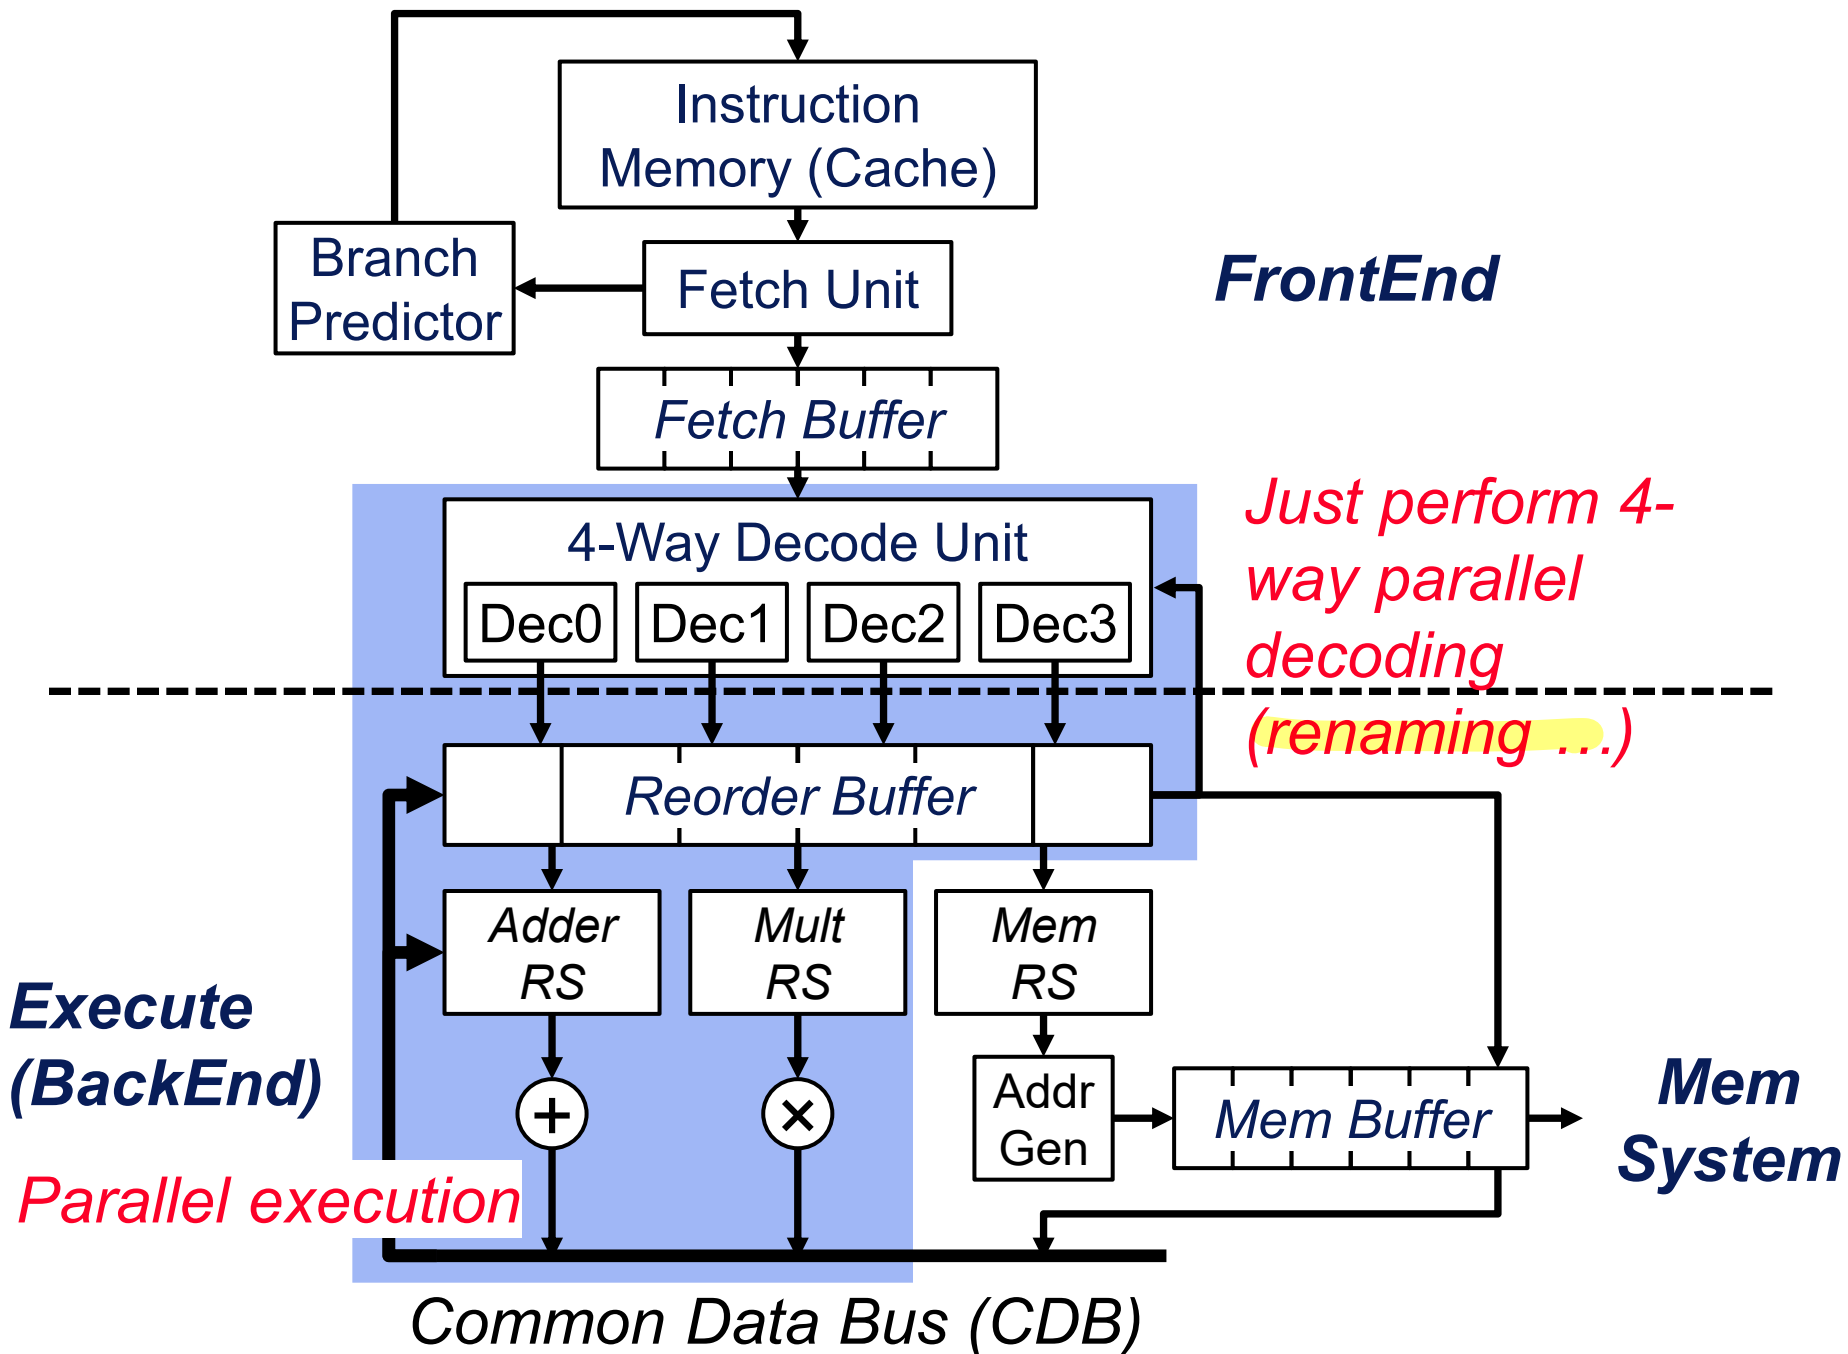


**Execute
(BackEnd)**

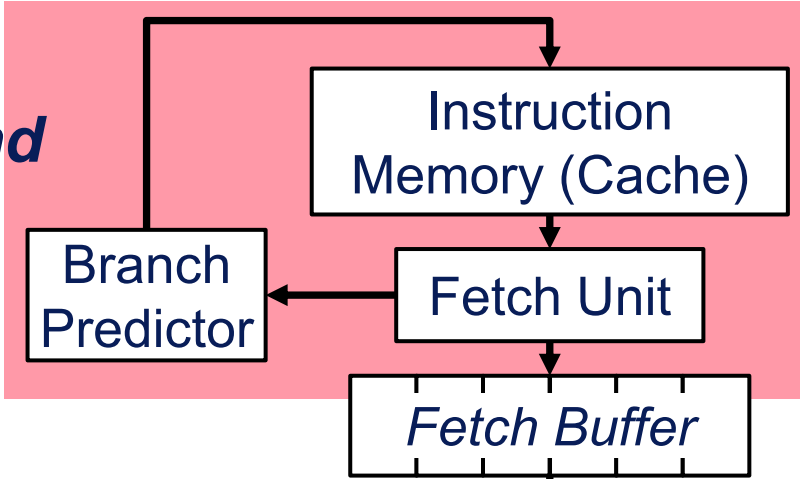


**Mem
System**

Common Data Bus (CDB)




FrontEnd



We need to simultaneously fetch multiple instructions!

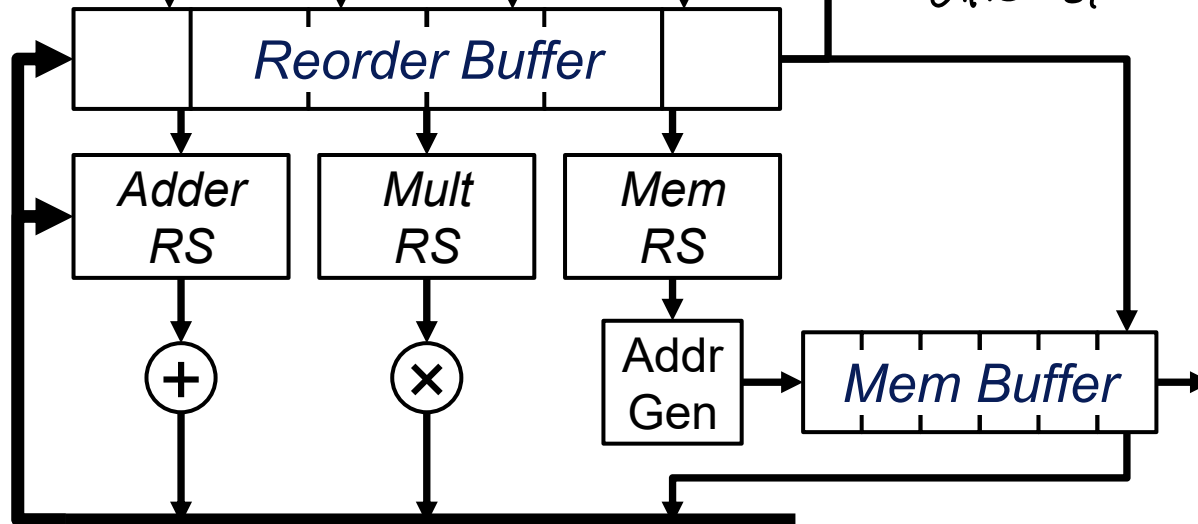
→ Fetch도 여러개 하자.



문제점

1. memory / code = (Alignment)
2. 동시에 Fetch된 것의 여의 Branch가 존재할 경우

Execute (BackEnd)



Mem System


Common Data Bus (CDB)

Fetching Multiple Instructions

- ◆ **Problem #1:** alignment problem with memory (cache)
 - You are not able to read instructions at different cache lines in parallel
- ◆ **Problem #2:** branches in-between N instruction sequence?
 - Should predict multiple branches in parallel
 - This incurs additional alignment issues

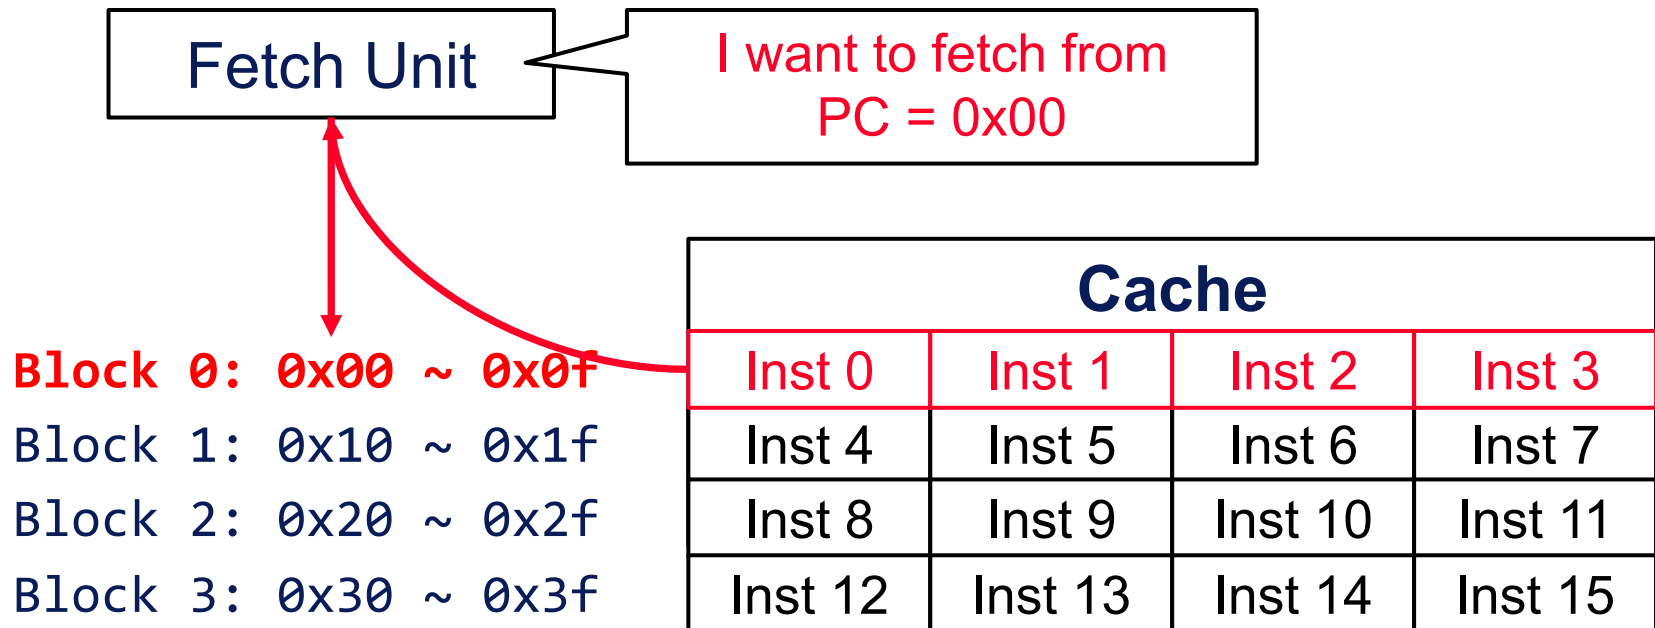
Preview on the Cache

- Cache is designed to read multiple words within a single cache block (typically 4~8 words) *Single Block에 여러 words 읽을 수 있는 구조: Block*
- The CPU can access 4 ~ 8 instructions in parallel (only if they are within the same block) *Single Block 안에 4~8개의 word*



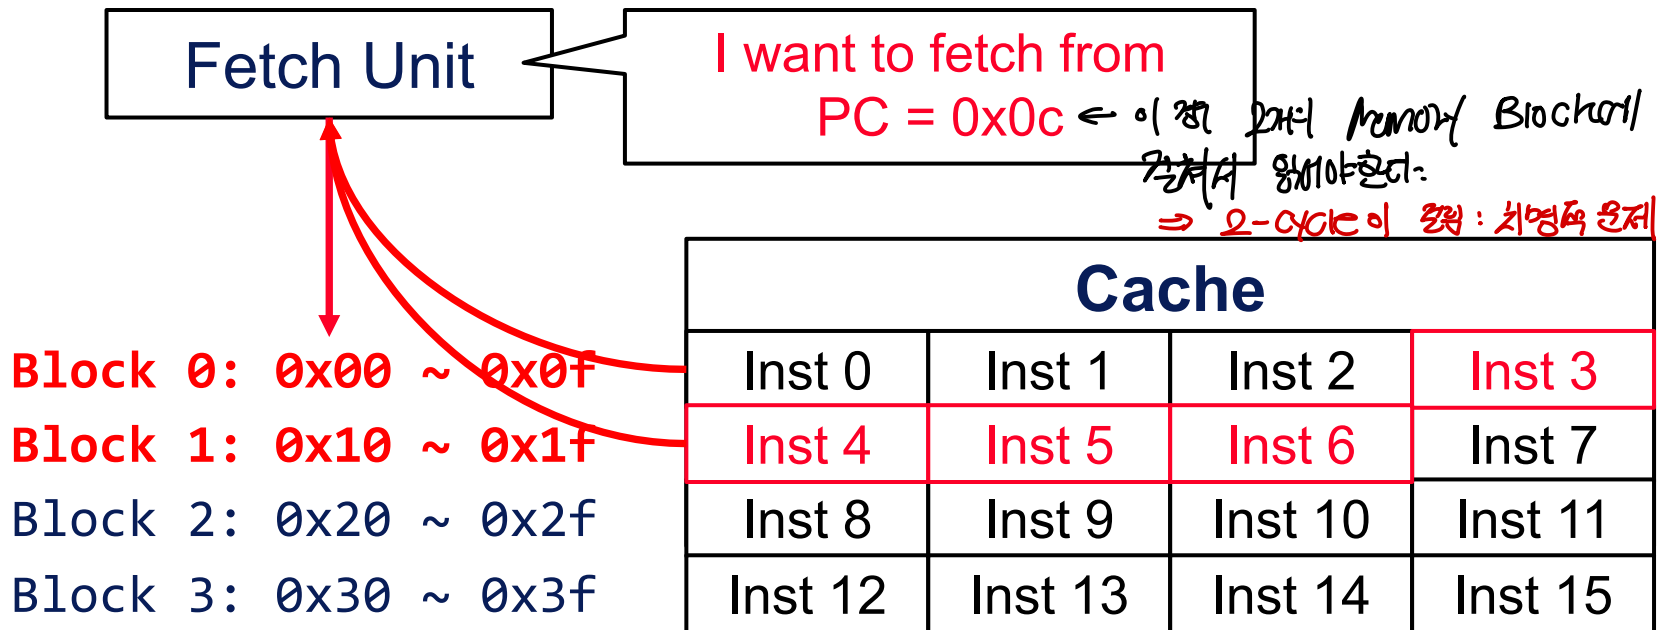
Cache misalignment problem

- ◆ Consider a cache with the block size of four words to support four-way superscalar



Cache misalignment problem

- ◆ Consider a cache with the block size of four words to support four-way superscalar



*There is one instruction @ Block 0
While others are @ Block 1*

→ Fetch becomes the bottleneck

Split line fetching

가장 쉬운 Solution

- ◆ You can divide a cache to keep even blocks and odd blocks separately → You can access the two blocks in parallel!

Fetch Unit

I want to fetch from
PC = 0x0c

Block 양이 많을 경우

2개

→ Cache Bank 2개를

각각 읽어들

→ Access ↑, Energy ↑

(2배 증가된 블록 대비)

Block 짝 Cache Bank 0			
Inst 0	Inst 1	Inst 2	Inst 3
Inst 8	Inst 9	Inst 10	Inst 11

Block 홀 Cache Bank 1			
Inst 4	Inst 5	Inst 6	Inst 7
Inst 12	Inst 13	Inst 14	Inst 15

Inst
3~6

Block 0: 0x00 ~ 0x0f

Block 2: 0x20 ~ 0x2f

Block 1: 0x10 ~ 0x1f

Block 3: 0x30 ~ 0x3f


Fetching Multiple Instructions

- ◆ **Problem #1:** alignment problem with memory (cache)
 - You are not able to read instructions at different cache lines in parallel

- ◆ **Problem #2:** branches in-between N instruction sequence?

한개 Instruction 중 여러 개의 Branch가 있을 수 있다

- Should predict multiple branches in parallel
- This incurs additional alignment issues



1. 여러 개의 Branch를 parallel하게 예측할 수 있어야 한다
2. Branch ↑ → Memory Alignment 문제 ↑

Increasing Branch Throughput

- ◆ There are a lot of branch instructions (15 ~ 20% of the instructions are branches)
- ◆ It is likely that there can be multiple branch instructions if you fetch N instructions (for superscalar)


Benchmark	taken %	avg basic block size	# instr between taken branches
eqntott	86.2%	4.20	4.87
espresso	63.8%	4.24	6.65
xlisp	64.7%	4.34	6.70
gcc	67.6%	4.65	6.88
sc	70.2%	4.71	6.71
compress	60.9%	5.39	8.85

There is a branch every 4 ~ 8 instructions


Increasing Branch Throughput [ICS'93]

- ◆ Let's say that there are an average of four instructions per basic block
 - ➔ Then, the maximum superscalar width would be four
- ◆ Extend Gag to predict multiple branches (MGag)

Remember Gag?




Pattern History Table (PHT)



Increasing Branch Throughput [ICS'93]


- ◆ Let's say that there are an average of four instructions per basic block
 - ➔ Then, the maximum superscalar width would be four
- ◆ Extend Gag to predict multiple branches (MGag)



First branch: 101 (global history)


Second branch: 01X (global history)

*Pattern History
Table (PHT)*



Branch-induced fragmentation

- ◆ Cache is good at accessing **contiguous data** (a single block contains multiple contiguous words)
- ◆ Taken instructions are placed in a **non-contiguous** region
- ◆ However, there are tons of taken branches in between (around 10~20%)



Branch-induced fragmentation


- ◆ Cache is good at accessing **contiguous data** (a single block contains multiple contiguous words)
- ◆ Taken instructions are placed in a **non-contiguous** region
- ◆ However, there are tons of taken branches in between (around 10~20%)

Benchmark	taken %	avg basic block size	# instr between taken branches
eqntott	86.2%	4.20	4.87
espresso	63.8%	4.24	6.65
xlisp	64.7%	4.34	6.70
gcc	67.6%	4.65	6.88
sc	70.2%	4.71	6.71
compress	60.9%	5.39	8.85

Every 4~8 instructions are placed in non-contiguous region

Trace caching [MICRO'96]

- ◆ Trace cache allocates instructions from different basic blocks (contiguous instructions) in a single cache line




Cache			
A0	A1	A2	A3
A4	B0	B1	B2
B3	B4	-	-
-	-	-	-
-	C0	C1	-
-	-	-	-
D0	D1	D2	D3
-	-	-	E1
E2	E3	-	-

Trace caching [MICRO'96]

- ◆ Trace cache allocates instructions from different basic blocks (contiguous instructions) in a single cache line

Cache			
A0	A1	A2	A3
A4	B0	B1	B2
B3	B4	-	-
-	-	-	-
-	C0	C1	-
-	-	-	-
D0	D1	D2	D3
-	-	-	E1
E2	E3	-	-

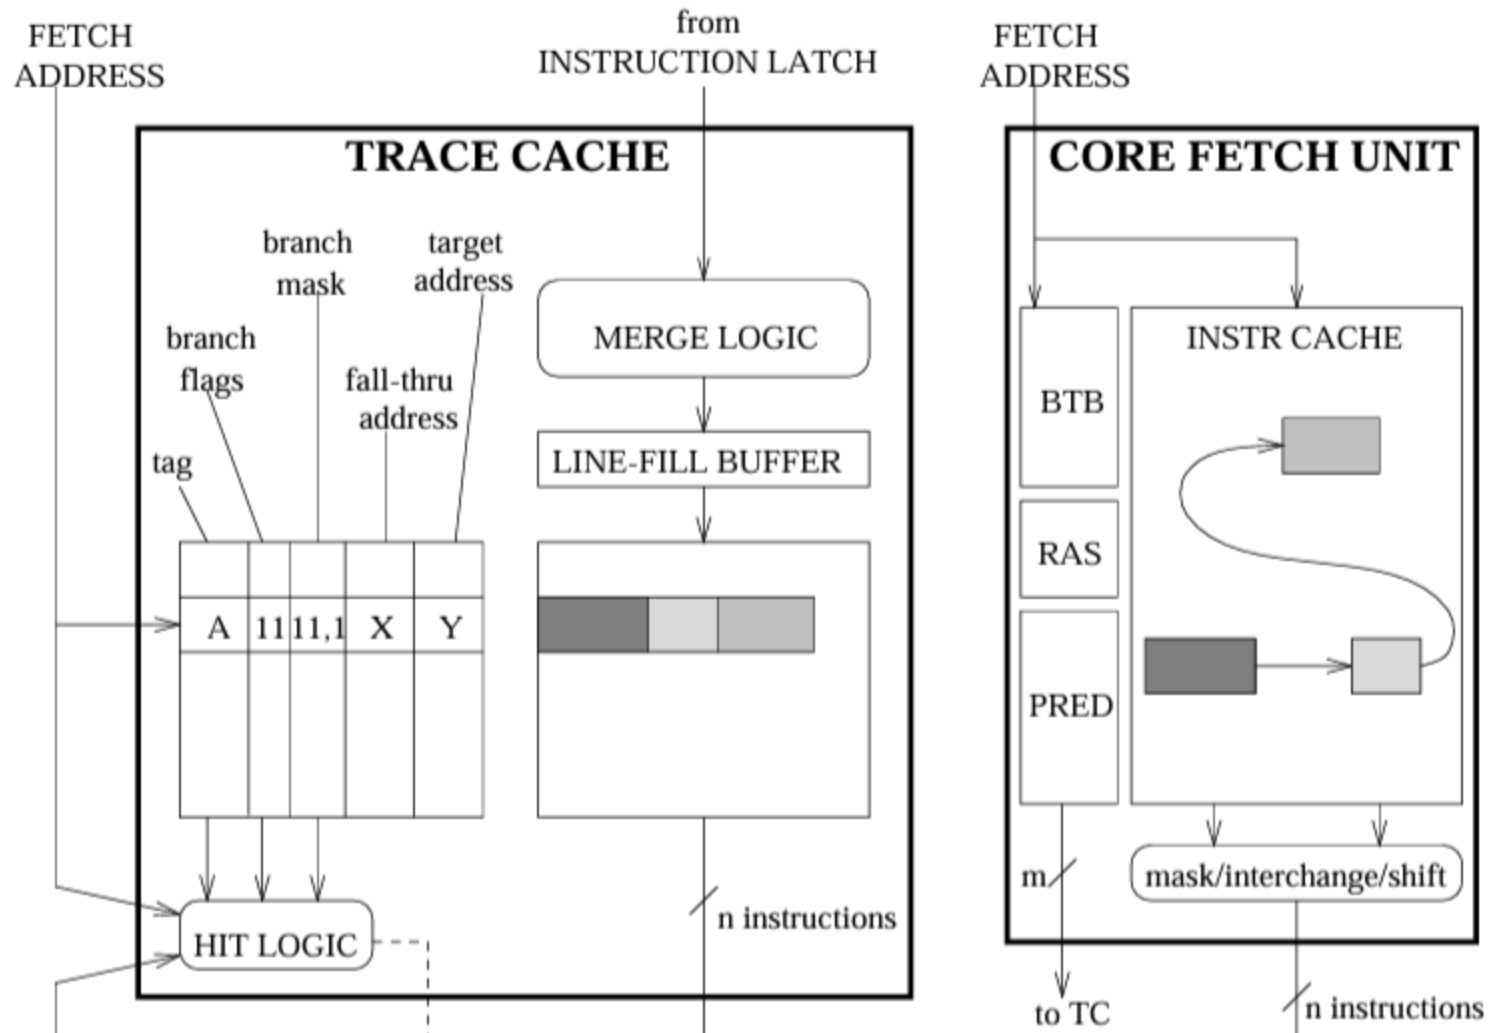


Trace Cache			
Metadata: A0, 0111, ...			
A0	A1	A2	A3
A4	B0	B1	B2
B3	B4	C0	C1
D1	D2	D3	D4
E1	E2	E3	-
-	-	-	-
-	-	-	-
-	-	-	-

*Use this data if
the predictions are
0 -> 1 -> 1 -> 1*

Trace caching

[MICRO'96]



Question?

Announcements:

Reading: finish reading P&H Ch.4

Handouts: none