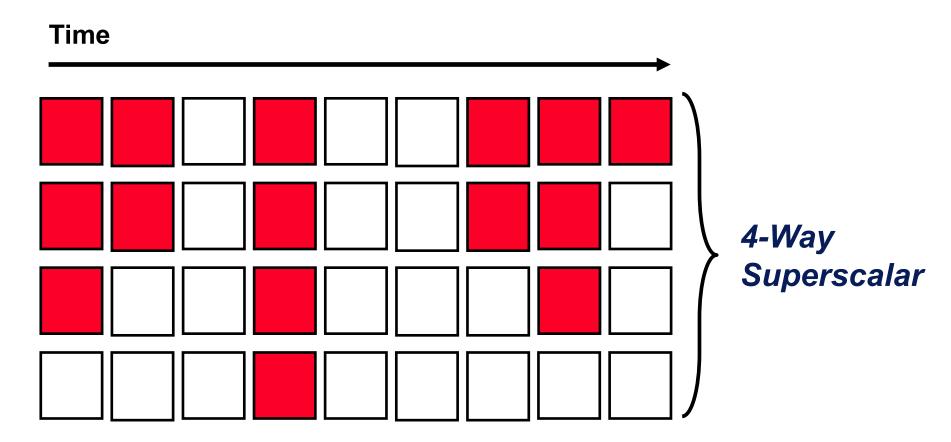
# Lecture 14: CPU – Multithreading

Hunjun Lee <a href="hunjunlee@hanyang.ac.kr">hunjunlee@hanyang.ac.kr</a>

## So far ...

- We learned about ways to minimize the idle cycles
  - Data forwarding
    - Use the data that are yet to be committed
  - Branch predictions
    - Predict the next instruction addresses for control instructions
  - Out-of-order execution
    - Execute instructions that do not suffer from true dependencies
  - Load-store queue
    - Prioritize load to resolve dependencies first (+ support load forwarding ...)
- This greatly reduces the idle cycles, but is this sufficient?
  - Deeper pipelines, long latency memory operations, limited ROB capacity and RS, Branch mispredictions, superscalar, ...

### Underutilization issues ...



It is extremely hard to fully utilize the entire pipeline due to large number of arithmetic units, superscalar, branch mispredictions, ......

## Performance and utilization

- ◆ Utilization → (actual IPC / peak IPC)
  - Even though the CPU has a large number of units, it cannot use them ...
- Even moderate superscalars (e.g., 4-way) are not fully utilized
  - IPC:  $1.5 \sim 2 \rightarrow < 50\%$  utilization
- Multithreading (MT)
  - One thread cannot fully utilize CPU (maybe we can with more threads)
  - Improve utilization by executing multiple threads on a single CPU

# **Basics of Multithreading**

- Thread basics
  - Instruction stream with state (registers and memory)
  - Register state is also called a thread context
- The threads can be both from the same process or from different processes
- You can either use software or hardware to support for multithreading!
  - SW: The CPU stores the context to the memory and load it back again
  - HW: The old thread's context is kept in a separate register file
    - What's mostly done in modern CPUs

## Latency vs. Throughput

- MT trades latency for throughput
  - Sharing processor degrades latency of individual threads
  - + But improves aggregate latency of both threads
  - + Improves utilization
- Example ...
  - Thread A: individual latency=10s, latency with thread B=15s
  - Thread B: individual latency=20s, latency with thread A=25s
    - Sequential latency (first A then B or vice versa): 30s
    - Parallel latency (A and B simultaneously): 25s
  - MT slows each thread by 5s
  - + But improves total latency by 5s

## How to run multiple threads?

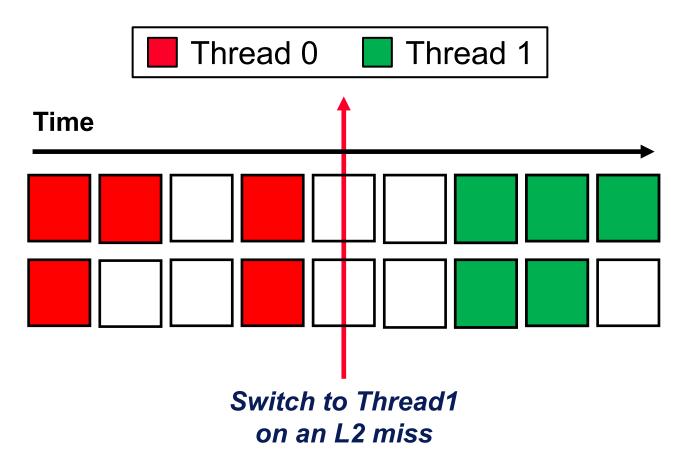
### Time sharing

- Execute a single thread at a time
- Switch threads on a long-latency operation (e.g., cache miss)
- Also known as "switch-on-miss" multithreading

### Space sharing

- Simple way (Fine-grain multithreading): execute different threads across pipeline depth (i.e., pipeline interleaving)
  - Fetch a different thread each cycle
- Advanced method (Simultaneous multithreading): both across pipeline depth + width
  - Fetch and issue each cycle from multiple threads
  - Need a policy to decide from which to thread fetch

## Coarse-grain Multithreading



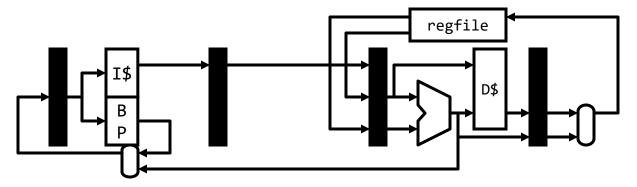
Multithreading minimizes single-thread performance but only hide very long latencies

# Coarse-Grain Multithreading (CGMT)

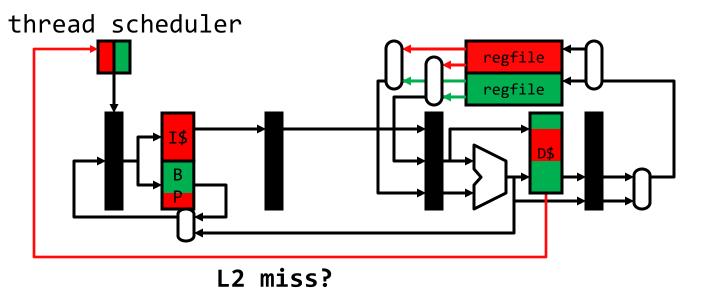
- Coarse-Grain Multi-Threading (CGMT)
  - Thread scheduling policy
    - Designate a "preferred" thread (e.g., thread A)
    - Switch to thread B on thread A's L2 miss
    - Switch back to A when A's L2 miss returns
  - + Sacrifices very little single thread performance (of one thread)
  - Tolerates only long latencies (e.g., L2 misses)
  - No pipeline partitioning
    - They flush on switch
      - Can't tolerate latencies shorter than twice pipeline depth
      - Why? We need two context switching → which incurs "at least" twich pipeline depth penalty
      - Need short in-order pipeline for good performance

## **CGMT**

### Normal Superscalar

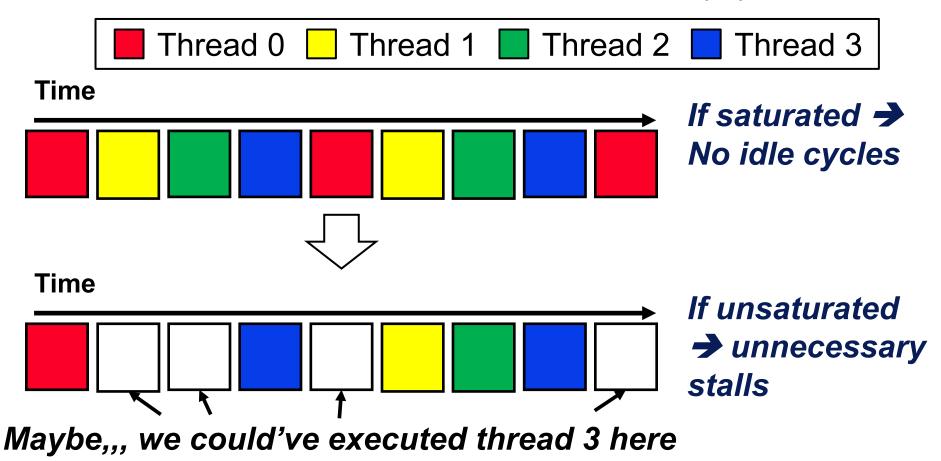


### CGMT



# Fine-grain Multithreading (FGMT)

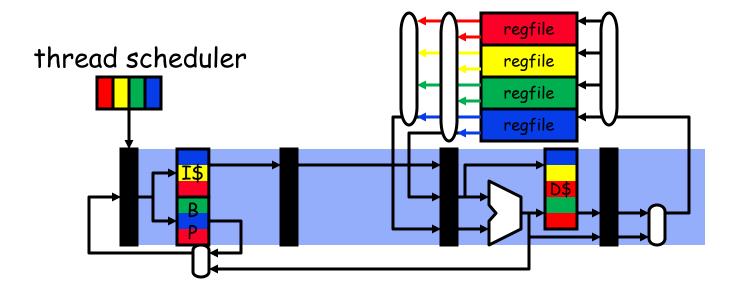
- Why not change the threads very frequently!
  - Extremely simple multithreading strategies
  - Fetch an instruction from different threads every cycle!



# Fine-Grain Multithreading

#### FGMT

- (Many) more threads
- Multiple threads in pipeline at once



# Illustrating FGMT: Resource View

Thread 0 Thread 1 Thread 2

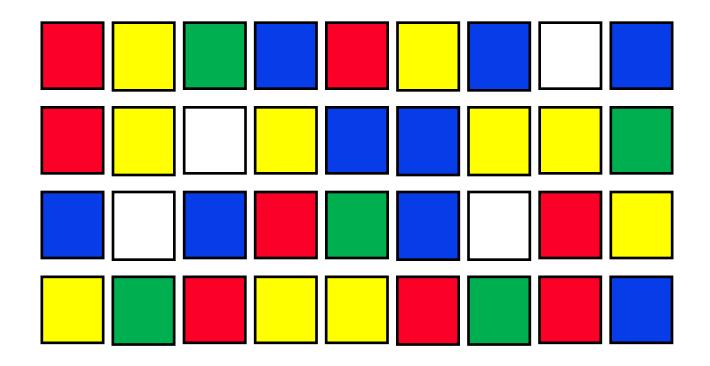
Thread 3 Thread 4

	t <sub>o</sub>	t <sub>1</sub>	t <sub>2</sub>	t <sub>3</sub>	t <sub>4</sub>	t <sub>5</sub>	t <sub>6</sub>	t <sub>7</sub>	t <sub>8</sub>	t <sub>9</sub>	t <sub>10</sub>
IF	I <sub>0</sub>	I <sub>1</sub>	<b>I</b> <sub>2</sub>								
ID		I <sub>0</sub>	I <sub>1</sub>								
EX			I <sub>0</sub>	I <sub>1</sub>	I <sub>1</sub>	I <sub>1</sub>	I <sub>1</sub>				
MEM				I <sub>0</sub>	I <sub>1</sub>	l <sub>1</sub>	I <sub>1</sub>				
WB					I <sub>0</sub>	I <sub>1</sub>	I <sub>1</sub>				

# Fine-grain Multithreading (FGMT)

- Although FGMT potentially suffers from unnecessary stalls, it greatly simplifies the CPU
- Example: Let's say the CPU is pipelined into 10 stages and supports 10-thread FGMT
  - Only a single instruction for each thread exists in the entire pipeline
    - No forwarding support
    - No branch prediction
    - No OoO support
    - No flush on switch
    - etc
- It is not widely adopted today

# Simultaneous Multithreading (SMT) (a.k.a. Hyperthreading)

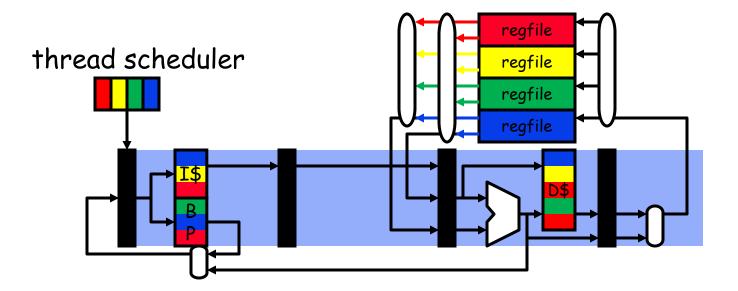


SMT maximizes the utilization (dispatch independent instructions which may or may not be from the same thread)

# Fine-Grain Multithreading

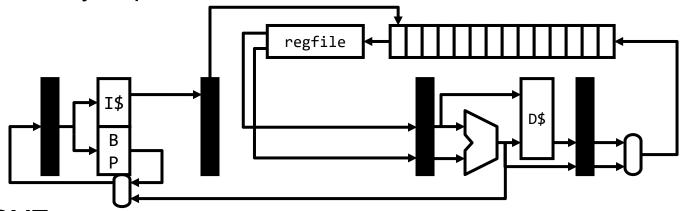
#### FGMT

- (Many) more threads
- Multiple threads in pipeline at once

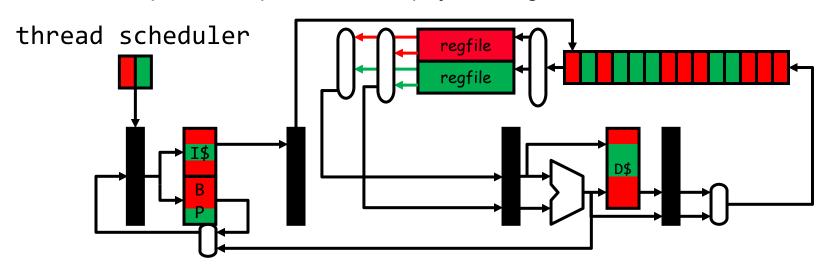


# Simultaneous Multithreading (SMT)

Ordinary Superscalar



- SMT
  - Replicate map table, share physical register file



# Simultaneous Multithreading (SMT) (a.k.a. Hyperthreading)

- Can we multithread an out-of-order machine?
  - Don't want to give up performance benefits
  - Don't want to give up natural tolerance of long memory latency
- Simultaneous multithreading (SMT)
  - + Tolerates all latencies (e.g., L2 misses, mispredicted branches)
  - ± Sacrifices some single thread performance due to resource sharing
  - Thread scheduling policy
    - Round-robin (just like FGMT) or something else?
  - Pipeline partitioning
    - Dynamic, but hard to be optimal
  - Example:
    - Pentium4 (hyper-threading): 5-way issue, 2 threads ...
    - Intel Xeon processors

### Potential issues in SMT

### Fetch policies

- Similar to FGMT, switch PCs every cycle
- Or flexible independent fetches (e.g., 2.8 fetch policy)
  - Can fetch from two threads (up to eight instructions), but fetch different number of instructions depending on the stalls

#### Cache interference

- General concern for all MT variants
- Can the working sets of multiple threads fit in the caches?
- Memory-sharing thread applications can help here
  - + Same instructions → share I\$,
  - + Shared data → less D\$ contention
  - MT is good for "server" workloads (e.g., database, web server)
- Large map table and physical register file
  - #phys-regs = (#threads \* #arch-regs) + #in-flight insns

## Potential issues in SMT

- What about branch predictor?
  - Without SMT: a single thread fills up the branch predictor
     (PHT) + branch target buffer
  - Utilizes the history to index the PHT and makes a predictor
     → This is good
  - How about SMT?
    - Share global history buffer?
    - Share pattern history table?
    - Separate branch target buffer?

## Potential issues in SMT

- How are ROB/LSQ/RS (or cache) partitioned in SMT?
  - There are two ways: Static and Dynamic

### Static partitioning

- Divide the hardware into equal-sized partitions
- Ensures fairness among different threads (some threads may occupy the slots ...)
- But, suffers from low utilization

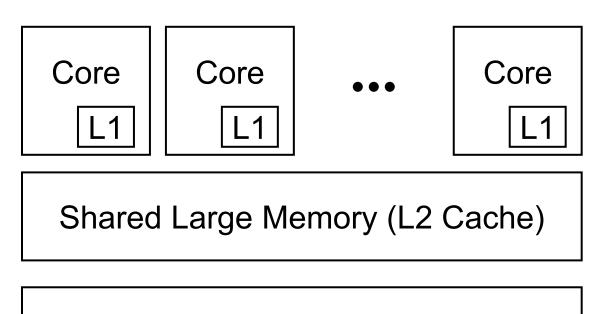
### Dynamic partitioning

- Dynamically resize the partition among different threads
- Higher utilization
- But, suffers from fairness issues
- Which thread to fetch? → e.g., fetch a thread having the fewest inflight instructions

There is no one-for-all solution (e.g., AMD → dynamic cache partition, Intel → static cache partition)

### Other method

- SMT is designed to execute multiple threads within a single core ...
- Alternatively, we can use multiple cores instead ...



Shared Larger Memory (L3 Cache)

## SMT vs. CMP

- If you wanted to run multiple threads would you build
  - Chip multiprocessor (CMP): multiple separate pipelines?
  - A multithreaded processor (SMT): a single larger pipeline?
- Both will get you throughput on multiple threads
  - CMP will be simpler, possibly faster clock
  - SMT will get you better performance (IPC) on a single thread
- Actually modern CPUs use both of them ...
  - Mostly, hyperthreading (2 threads per core) \* multiple cores
- But, more recently, Intel removed Hyper-threading ...
  - Why? The modern workloads are bounded by the memory bandwidth (not latency) → Think about AI services and databases ...

# Sidenote: Defeaturing things

- Computer architects are developing tons of new features for higher performance
- We are afraid of "de"featuring what they have already developed!
- But, sometimes, removing features can be as important as adding a new feature (especially as the architects are becoming more and more complex)

## Question?

Announcements:

Reading: finish reading P&H Ch.6.4

Handouts: none